Summer Internship Presentation

Nikhil Shekhar
Navdeep Gill
08/12/2016

Assignment Details:

Develop a framework that can help benchmark multiple machine learning libraries

Background

Education

- Masters in Computer Science at State university of New York at Buffalo - **Graduating December 2016**.
- Bachelor of Engineering in Computer Science.

Work Experience

- 4 years in Big Data Eco-space in companies - Impetus technologies, Center of Excellence at Cognizant.
- Worked on multiple Big data projects during the 4 years for clients like Capital One, American express, AT&T, Cisco, Lexus.
- 1 year in Java - J2EE technologies.

Project details

ML Frameworks used for benchmarking efforts:

- Spark-ml
- Xgboost
- scikit

Benchmark workflow

- Read the test case file and get the Algorithm details and the hyper-parameters for the same.
- Connect to EC2 S3 storage and read the data files needed for the test case.
- Generate the feature matrix
- Run the algorithm
- Collect the accuracy metrics
- Collect the performance metrics - training and scoring time
- Push the metrics to MySQL for each test case and each run
- Logs are written to a configurable directory

contd...

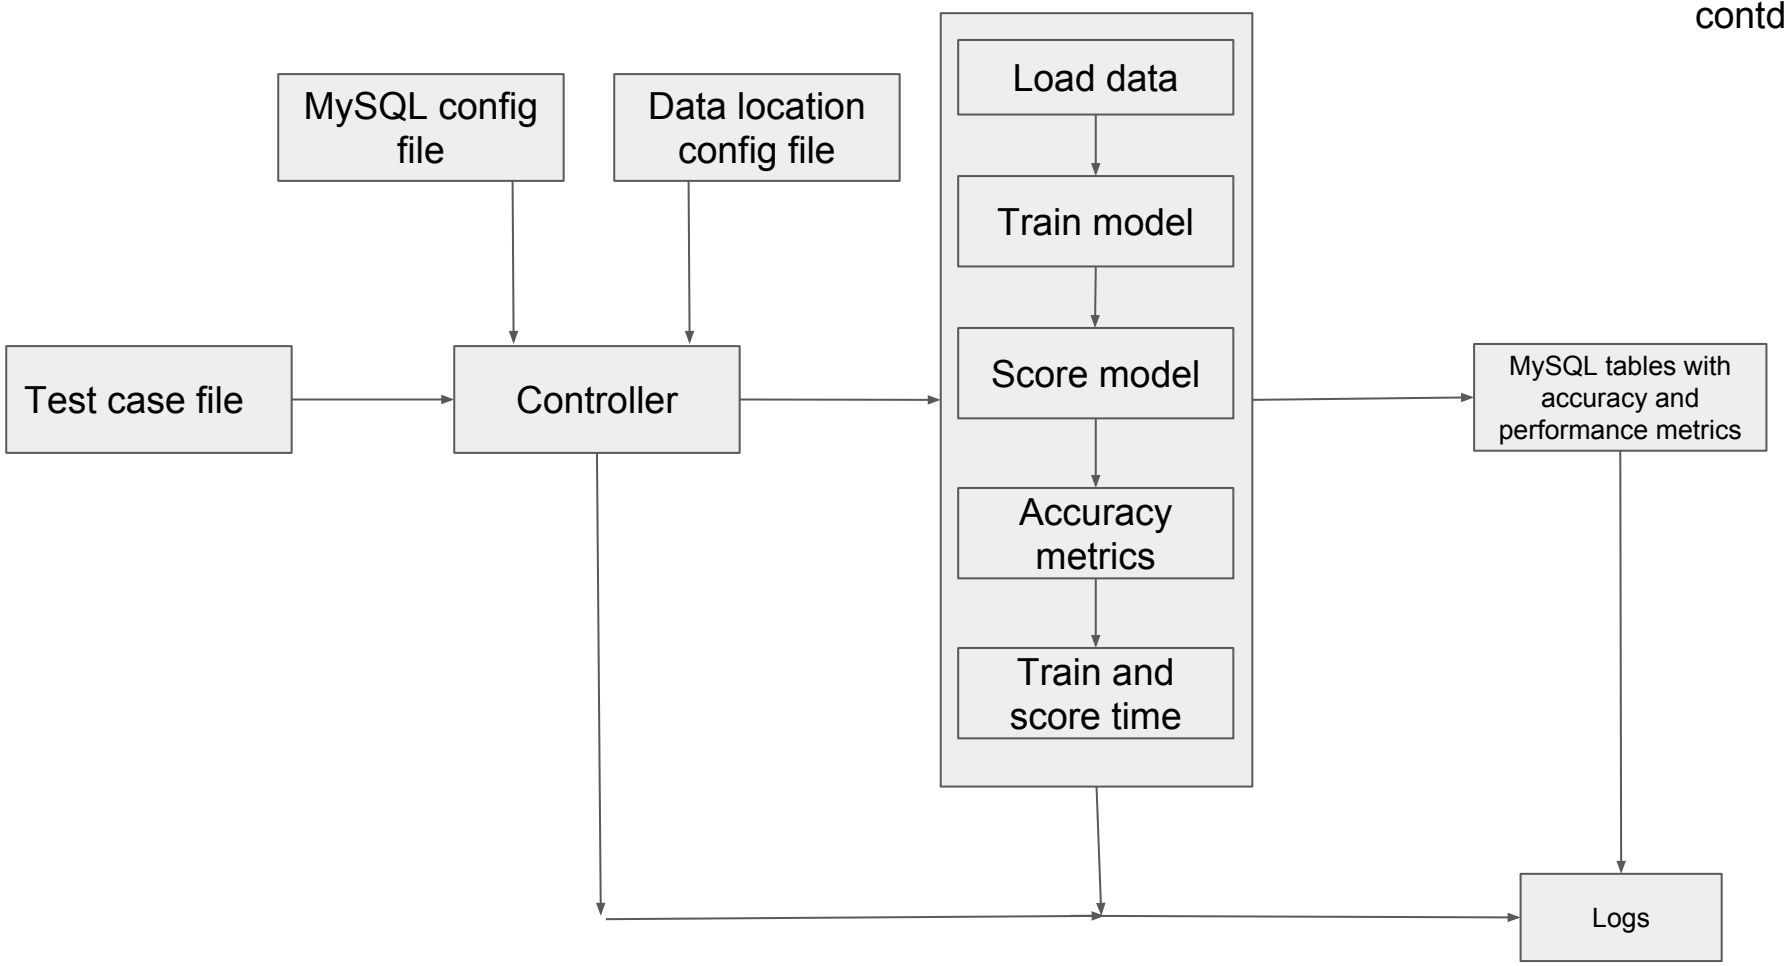
MySQL config file
Data location config file
Load data
Train model
Score model
Accuracy metrics
Train and score time
Test case file
Controller
MySQL tables with accuracy and performance metrics
Logs

Preprocessing data for xgboost and scikit

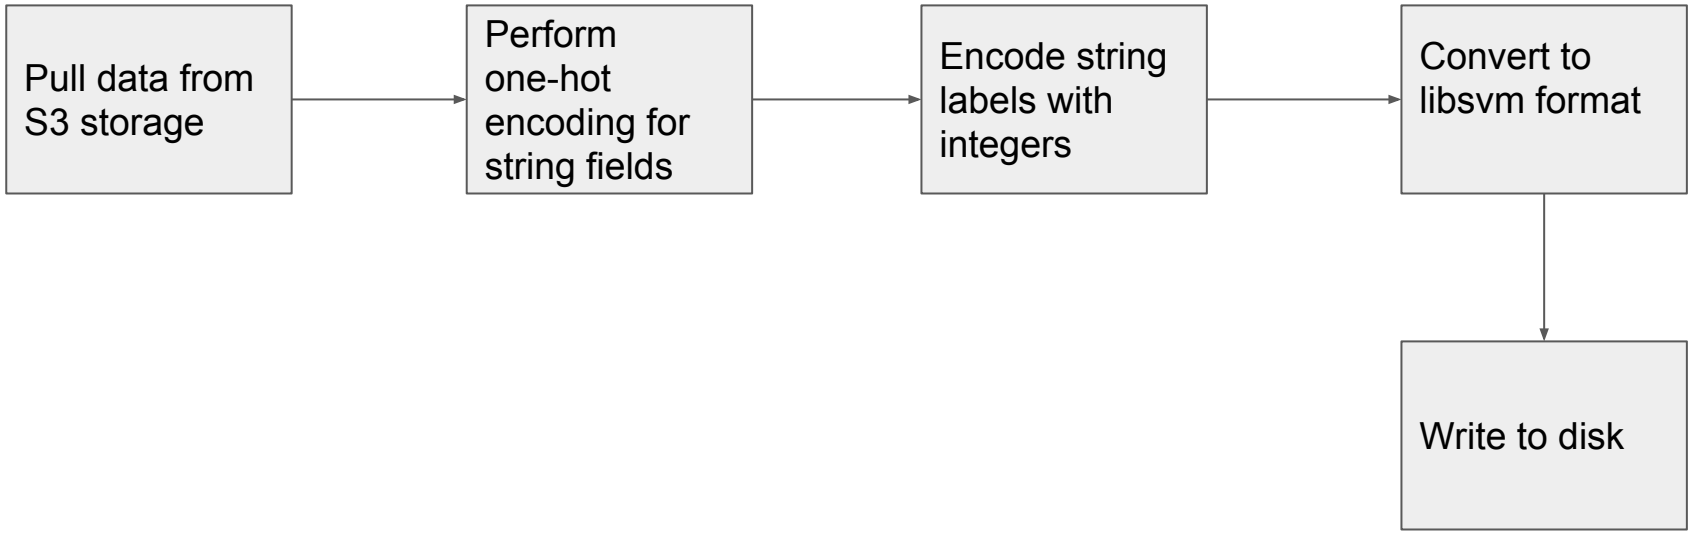

Benchmark - MLlib

Algorithms supported:

- Linear regression
- Logistic regression - Supports only binary classification.
- Random forest classifier
- Random forest regressor
- Gradient boosted machine classifier - Supports only binary classification.
- Gradient boosted machine regressor
- Perceptron classifier
- KMeans
- Principal component analysis

Achieved through Java API’s on Spark - 1.6.2

Benchmark - xgboost

Algorithms supported:

- Gradient Boosting algorithm that supports regression and classification.

- A pre-processing step is needed which pulls data from S3, performs one-hot encoding and creates a libsvm file which is fed to xgboost
- Achieved through Java API’s on xgboost - 0.6

Benchmark - scikit learn

Algorithms supported:

- Linear regression
- Logistic regression
- Random forest classifier
- Random forest regressor
- Gradient boosted machine classifier
- Gradient boosted machine regressor
- Naive Bayes
- KMeans

- A pre-processing step is needed which pulls data from S3 and performs one-hot encoding
- Achieved through Python API’s on scikit - 0.17

Test cases file

```
test_case_id,algorithm,training_data_file_id,test_data_file_id,parameters,grid_parameters,nfold,multiclass(T/F),header(Y/N),number_of_colmns_if_no_header  #10c
1,glm,5,6,param_regularizer:0.3;param_elastic_net:0.8;param_number_of_iteration:10,param_regularizer_grid:0.1:0.5;param_elastic_net_grid:0.1:1,nfold:0,,N,90,,
##2,glm,19,20,param_regularizer:0.3;param_elastic_net:0.8;param_number_of_iteration:10,param_regularizer_grid:0.1:0.5;param_elastic_net_grid:0.1:1,nfold:0,,Y,,
3,glm,21,22,param_regularizer:0.3;param_elastic_net:0.8;param_number_of_iteration:10,param_regularizer_grid:0.1:0.5;param_elastic_net_grid:0.1:1,nfold:0,,Y,,
4,glm,23,24,param_regularizer:0.3;param_elastic_net:0.8;param_number_of_iteration:10,param_regularizer_grid:0.1:0.5;param_elastic_net_grid:0.1:1,nfold:0,,Y,,
5,rfc,3,4,param_number_of_trees:3;param_max_depth:10,param_max_depth_grid:1:3;param_number_of_trees_grid:1:2,nfold:0,T,Y,,
##6,rfc,5,6,param_number_of_trees:3;param_max_depth:10,param_max_depth_grid:1:3;param_number_of_trees_grid:1:2,nfold:0,N,N,90,
7,rfc,7,8,param_number_of_trees:3;param_max_depth:3;param_max_bins:110,param_max_depth_grid:1:3;param_number_of_trees_grid:1:2,nfold:0,T,Y,,
8,rfc,9,10,param_number_of_trees:3;param_max_depth:10,param_max_depth_grid:1:3;param_number_of_trees_grid:1:2,nfold:0,F,Y,,
##9,rfc,11,12,param_number_of_trees:3;param_max_depth:10,param_max_depth_grid:1:3;param_number_of_trees_grid:1:2,nfold:0,T,Y,55,
10,rfc,13,14,param_number_of_trees:3;param_max_depth:10,param_max_depth_grid:1:3;param_number_of_trees_grid:1:2,nfold:0,T,Y,785,
##15,rfc,1,2,param_number_of_trees:3;param_max_depth:10,param_max_depth_grid:1:3;param_number_of_trees_grid:1:2,nfold:0,T,Y,785,
##17,rfc,15,16,param_number_of_trees:3;param_max_depth:10,param_max_depth_grid:1:3;param_number_of_trees_grid:1:2,nfold:0,T,Y,785,
##18,rfc,17,18,param_number_of_trees:3;param_max_depth:10,param_max_depth_grid:1:3;param_number_of_trees_grid:1:2,nfold:0,T,Y,785,
19,rfr,5,6,param_number_of_trees:3;param_max_depth:3,param_max_depth_grid:1:3;param_number_of_trees_grid:1:2,nfold:0,N,N,90,
##20,rfr,19,20,param_number_of_trees:3;param_max_depth:3,param_max_depth_grid:1:3;param_number_of_trees_grid:1:2,nfold:0,N,Y,,
21,rfr,21,22,param_number_of_trees:3;param_max_depth:3,param_max_depth_grid:1:3;param_number_of_trees_grid:1:2,nfold:0,N,Y,,
22,rfr,23,24,param_number_of_trees:3;param_max_depth:3,param_max_depth_grid:1:3;param_number_of_trees_grid:1:2,nfold:0,N,Y,,
##23,gbmr,5,6,param_number_of_iteration:2;param_max_depth:2,param_max_depth_grid:1:3,nfold:0,N,N,90,
##24,gbmr,19,20,param_number_of_iteration:2;param_max_depth:2,param_max_depth_grid:1:3,nfold:0,N,Y,,
25,gbmr,21,22,param_number_of_iteration:2;param_max_depth:2,param_max_depth_grid:1:3,nfold:0,N,Y,,
26,gbmr,23,24,param_number_of_iteration:2;param_max_depth:2,param_max_depth_grid:1:3,nfold:0,N,Y,,
27,gbmc,3,4,param_number_of_iteration:2;param_max_depth:2,param_max_depth_grid:1:3,nfold:0,F,Y,,
28,gbmc,7,8,param_number_of_iteration:2;param_max_depth:2;param_max_bins:110,param_max_depth_grid:1:3,nfold:0,F,Y,,
29,gbmc,9,10,param_number_of_iteration:2;param_max_depth:2,param_max_depth_grid:1:3,nfold:0,F,Y,,
##32,gbmc,15,16,param_number_of_iteration:2;param_max_depth:2,param_max_depth_grid:1:3,nfold:0,T,Y,,
##33,gbmc,17,18,param_number_of_iteration:2;param_max_depth:2,param_max_depth_grid:1:3,nfold:0,F,Y,,
34,nn,3,4,param_number_of_iteration:2;param_max_depth:2;param_layers_and_units:7#3#2,param_max_depth_grid:1:3,nfold:0,T,Y,,
35,nn,7,8,param_number_of_iteration:2;param_max_depth:2;param_layers_and_units:8#3#2,param_max_depth_grid:1:3,nfold:0,T,Y,,
##36,nn,9,10,param_number_of_iteration:2;param_max_depth:2;param_layers_and_units:10000#500#2,param_max_depth_grid:1:3,nfold:0,T,Y,,
37,nn,11,12,param_number_of_iteration:2;param_max_depth:2;param_layers_and_units:54#27#7,param_max_depth_grid:1:3,nfold:0,T,Y,55,
38,nn,14,14,param_number_of_iteration:2;param_max_depth:2;param_layers_and_units:784#350#10,param_max_depth_grid:1:3,nfold:0,T,Y,785,
##39,nn,15,16,param_number_of_iteration:2;param_max_depth:2;param_layers_and_units:7#3#2,param_max_depth_grid:1:3,nfold:0,T,Y,,
##40,nn,17,18,param_number_of_iteration:2;param_max_depth:2;param_layers_and_units:7#3#2,param_max_depth_grid:1:3,nfold:0,T,Y,,
41,glmlogistic,3,4,param_regularizer:0.3;param_elastic_net:0.8;param_number_of_iteration:10,param_regularizer_grid:0.1:0.5;param_elastic_net_grid:0.1:1,nfold:0
42,glmlogistic,7,8,param_regularizer:0.3;param_elastic_net:0.8;param_number_of_iteration:10,param_regularizer_grid:0.1:0.5;param_elastic_net_grid:0.1:1,nfold:0
##43,glmlogistic,10,10,param_regularizer:0.3;param_elastic_net:0.8;param_number_of_iteration:10,param_regularizer_grid:0.1:0.5;param_elastic_net_grid:0.1:1,nfo
##44,glmlogistic,17,18,param_regularizer:0.3;param_elastic_net:0.8;param_number_of_iteration:10,param_regularizer_grid:0.1:0.5;param_elastic_net_grid:0.1:1,nfo
##45,kmeans,9,10,param_number_of_iteration:2;param_number_of_clusters:5,,
#TEST CASES WITH CROSS VALIDATION
#GBT supports only binary classification
#Logistic regression supports only binary classification
```

Data file details

```
id,data_path,feature_matrix_columns,output_column_name,feature_matrix_column_if_string,label_column_name
1,s3n://h2o-public-test-data/smalldata/testng/iris_train1.csv,Sepal.Length;Sepal.Width;Petal.Length,features,,Petal.Width
2,s3n://h2o-public-test-data/smalldata/testng/iris_validation1.csv
3,s3n://h2o-public-test-data/smalldata/testng/prostate_train.csv,AGE;RACE;DPROS;DCAPS;PSA;VOL;GLEASON,features,,CAPSULE
4,s3n://h2o-public-test-data/smalldata/testng/prostate_test.csv
5,s3n://h2o-public-test-data/bigdata/laptop/testng/milsongs-train.csv.gz,,test,,,
6,s3n://h2o-public-test-data/bigdata/laptop/testng/milsongs-test.csv.gz
7,s3n://h2o-public-test-data/smalldata/testng/airlines_train.csv,Distance,features,fYear;fMonth;fDayofMonth;fDayOfWeek;UniqueCarrier;Origin;Dest,IsDepDelayed
8,s3n://h2o-public-test-data/smalldata/testng/airlines_test.csv
9,s3n://h2o-public-test-data/smalldata/testng/arcene_train.csv,X1;C2;C3;C4;C5;C6;C7;C8;C9;C10;C11;C12;C13;C14;C15;C16;C17;C18;C19;C20;C21;C22;C23;C24;C25;C26;C27;C28;C29;C30;C31
10,s3n://h2o-public-test-data/smalldata/testng/arcene_test.csv,X1;C2;C3;C4;C5;C6;C7;C8;C9;C10;C11;C12;C13;C14;C15;C16;C17;C18;C19;C20;C21;C22;C23;C24;C25;C26;C27;C28;C29;C30;C31
11,s3n://h2o-public-test-data/smalldata/testng/covtype_small_dense_multiclass_unbalanced_train.csv
12,s3n://h2o-public-test-data/smalldata/testng/covtype_small_dense_multiclass_unbalanced_test.csv
13,s3n://h2o-public-test-data/bigdata/laptop/testng/mnist_train.csv
14,s3n://h2o-public-test-data/bigdata/laptop/testng/mnist_test.csv
15,s3n://h2o-public-test-data/bigdata/laptop/testng/cup98_train.csv
16,s3n://h2o-public-test-data/bigdata/laptop/testng/cup98_test.csv
17,s3n://h2o-public-test-data/bigdata/laptop/testng/higgs_train_imbalance_100k.csv
18,s3n://h2o-public-test-data/bigdata/laptop/testng/higgs_test_imbalance_100k.csv
19,s3n://h2o-public-test-data/smalldata/testng/cars_train.csv,economy (mpg);cylinders;power (hp);displacement (cc);weight (lb);year,features,name,0-60 mph (s)
20,s3n://h2o-public-test-data/smalldata/testng/cars_test.csv
21,s3n://h2o-public-test-data/smalldata/testng/housing_train.csv,C1;C2;C3;C4;C5;C6;C7;C8;C9;C10;C11;C12;C13,features,,C14
22,s3n://h2o-public-test-data/smalldata/testng/housing_test.csv
23,s3n://h2o-public-test-data/smalldata/testng/computer_train.csv,C3;C4;C5;C6;C7;C8,features,,C10
24,s3n://h2o-public-test-data/smalldata/testng/computer_test.csv
25,s3n://h2o-public-test-data/smalldata/testng/iris.csv
26,s3n://h2o-public-test-data/smalldata/testng/prostate.csv
27,s3n://h2o-public-test-data/bigdata/laptop/testng/milsongs.csv.gz
28,s3n://h2o-public-test-data/smalldata/testng/airlines.csv
29,s3n://h2o-public-test-data/smalldata/testng/arcene.csv
30,s3n://h2o-public-test-data/smalldata/testng/covtype_small_dense_multiclass_unbalanced.csv
31,s3n://h2o-public-test-data/bigdata/laptop/testng/mnist.csv
32,s3n://h2o-public-test-data/bigdata/laptop/testng/cup98.csv
33,s3n://h2o-public-test-data/bigdata/laptop/testng/higgs_imbalance_100k.csv
34,s3n://h2o-public-test-data/smalldata/testng/cars.csv
35,s3n://h2o-public-test-data/smalldata/testng/housing.csv
36,s3n://h2o-public-test-data/smalldata/testng/computer.csv
```

MySQL config

```
driver = com.mysql.jdbc.Driver
db = h2o
host = 172.16.2.178
user = root
password = 0xdata
table = SparkStats
port = 3306
```

MySQL metrics table

mysql> select * from SparkStats;

id	mae	r2	rmse	traintime	testtime	timestamp	dataset	algo	f1	precision	recall	weightedprecision	weightedrecall	wssse	auc	aupr	type
3	48.75362221547697	0.06964307748519938	59.163828296530475	1682	49	1470935165863	s3n://h2o-public-test-data/smalldata/testng/housing_train.csv	glm	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	Regression
4	14.674685374001873	0.43884691396525277	19.04954819588868	588	31	1470935173114	s3n://h2o-public-test-data/smalldata/testng/computer_train.csv	glm	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	Regression
5	NULL	NULL	NULL	3373	52	1470935180184	s3n://h2o-public-test-data/smalldata/testng/prostate_train.csv	rfc	0.8778758503027191	0.8786885245901639	0.8786885245901639	0.8786682672468771	0.8786885245901639	NULL	NULL	NULL	Classification
1	7.330103901417214	0.19146711437813813	10.061443158888475	21698	84	1470936439069	s3n://h2o-public-test-data/bigdata/laptop/testng/milsongs-train.csv.gz	glm	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	Regression
7	NULL	NULL	NULL	9127	57	1470936605295	s3n://h2o-public-test-data/smalldata/testng/airlines_train.csv	rfc	0.5665018106592777	0.6105810572867614	0.6105810572867614	0.6315342765569345	0.6105810572867614	NULL	NULL	NULL	Classification
8	NULL	NULL	NULL	63442	4528	1470936761732	s3n://h2o-public-test-data/smalldata/testng/arcene_train.csv	rfc	NULL	NULL	NULL	NULL	NULL	NULL	0.9683441558441559	0.9684659090909091	Classification
10	NULL	NULL	NULL	394712	428	1470937476723	s3n://h2o-public-test-data/bigdata/laptop/testng/mnist_train.csv	rfc	0.8999811145775389	0.90005	0.90005	0.9003406791716891	0.90005	NULL	NULL	NULL	Classification
19	7.883419419339239	0.10409521176058234	10.591131554352721	19153	61	1470940714924	s3n://h2o-public-test-data/bigdata/laptop/testng/milsongs-train.csv.gz	rfr	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	Regression
21	45.67486579111672	0.16188355038473445	56.15438219928148	265	28	1470940901104	s3n://h2o-public-test-data/smalldata/testng/housing_train.csv	rfr	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	Regression
22	13.447535390062225	0.5621724523771108	16.826567953132106	219	27	1470940917420	s3n://h2o-public-test-data/smalldata/testng/computer_train.csv	rfr	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	Regression
25	35.71034985861266	0.47059963360347745	44.62966120882825	15299	28	1470940921737	s3n://h2o-public-test-data/smalldata/testng/housing_train.csv	gbmr	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	Regression
26	7.132483202426509	0.8174783845356943	10.864285104581114	16513	27	1470940960023	s3n://h2o-public-test-data/smalldata/testng/computer_train.csv	gbmr	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	Regression
27	NULL	NULL	NULL	32174	23	1470940984230	s3n://h2o-public-test-data/smalldata/testng/prostate_train.csv	gbmc	NULL	NULL	NULL	NULL	NULL	NULL	0.8439449296025664	0.8561435414990957	Classification
28	NULL	NULL	NULL	8420	25	1470941048495	s3n://h2o-public-test-data/smalldata/testng/airlines_train.csv	gbmc	NULL	NULL	NULL	NULL	NULL	NULL	0.6669995513119844	0.7160024170166069	Classification
29	NULL	NULL	NULL	55948	2082	1470941143105	s3n://h2o-public-test-data/smalldata/testng/arcene_train.csv	gbmc	NULL	NULL	NULL	NULL	NULL	NULL	1.0	1.0	Classification
35	NULL	NULL	NULL	31386	18	1470941275577	s3n://h2o-public-test-data/smalldata/testng/airlines_train.csv	nn	0.3906508272825593	0.5471520412759511	0.5471520412759511	0.541831992988346	0.547152041275951	NULL	NULL	NULL	Classification
37	NULL	NULL	NULL	7973	71	1470942213971	s3n://h2o-public-test-data/smalldata/testng/covtype_small_dense_multiclass_unbalanced_train.csv	nn	0.7578330869011121	0.8335164835164836	0.8335164835164836	0.6947497282936844	0.8335164835164836	NULL	NULL	NULL	Classification
38	NULL	NULL	NULL	89785	218	1470942289940	s3n://h2o-public-test-data/bigdata/laptop/testng/mnist_test.csv	nn	0.5454015569671474	0.5912	0.5912	0.6581570379588935	0.5912000000000001	NULL	NULL	NULL	Classification
41	NULL	NULL	NULL	359	34	1470942567576	s3n://h2o-public-test-data/smalldata/testng/prostate_train.csv	glmlogistic	0.7809659597219747	NULL	NULL	NULL	NULL	NULL	NULL	0.7293973458371062	Classification
42	NULL	NULL	NULL	574	33	1470942584334	s3n://h2o-public-test-data/smalldata/testng/airlines_train.csv	glmlogistic	0.5	NULL	NULL	NULL	NULL	NULL	NULL	0.7265672986364194	Classification
34	NULL	NULL	NULL	5944	56	1470944697246	s3n://h2o-public-test-data/smalldata/testng/prostate_train.csv	nn	0.44202927882345	0.5934426229508196	0.5934426229508196	0.35217414673474867	0.5934426229508196	NULL	NULL	NULL	Classification

21 rows in set (0.03 sec)

contd..

```
mysql> select * from XgboostStats;
```

id	run	trainrmse	testrmse	trainmae	testmae	trainlogloss	testlogloss	trainrerror	testerror	trainauc	testauc	trainingtime	testtime	timestamp	dataset	algo	type
1	1	0.168564	0.170865	0.142931	0.144147	0.162225	0.164609	0.007318	0.007525	0.997955	0.997318	84	0	1470946864967	/Users/nikhilshekhar/h2o/xgboost/demo/data/agaricus.txt.train	gbm	gbm
1	2	0.075735	0.080230	0.053599	0.054797	0.059073	0.061924	0.002098	0.002918	0.999014	0.998326	84	0	1470946864967	/Users/nikhilshekhar/h2o/xgboost/demo/data/agaricus.txt.train	gbm	gbm
1	3	0.041020	0.047022	0.021570	0.022513	0.024808	0.027473	0.001024	0.001843	0.999065	0.998697	84	0	1470946864967	/Users/nikhilshekhar/h2o/xgboost/demo/data/agaricus.txt.train	gbm	gbm
1	4	0.027302	0.033732	0.009101	0.009886	0.011760	0.014158	0.000460	0.001382	0.999171	0.998816	84	0	1470946864967	/Users/nikhilshekhar/h2o/xgboost/demo/data/agaricus.txt.train	gbm	gbm
1	5	0.023020	0.027296	0.004178	0.004737	0.006097	0.007927	0.000460	0.000921	0.999806	0.999552	84	0	1470946864967	/Users/nikhilshekhar/h2o/xgboost/demo/data/agaricus.txt.train	gbm	gbm
2	1	0.416912	0.421947	0.360991	0.366655	0.522843	0.535534	0.263100	0.292763	0.818024	0.803852	15	0	1470946865906	/Users/nikhilshekhar/h2o_working_directory/test/s3data_new/prostate_train.lib	gbm	gbm
2	2	0.404101	0.412338	0.330628	0.336606	0.495632	0.521224	0.232533	0.233553	0.834429	0.818020	15	0	1470946865906	/Users/nikhilshekhar/h2o_working_directory/test/s3data_new/prostate_train.lib	gbm	gbm
2	3	0.402385	0.413338	0.326205	0.333154	0.491851	0.524826	0.222707	0.233553	0.833742	0.811651	15	0	1470946865906	/Users/nikhilshekhar/h2o_working_directory/test/s3data_new/prostate_train.lib	gbm	gbm
2	4	0.402228	0.413681	0.325030	0.332208	0.491452	0.525749	0.225982	0.240132	0.834279	0.811779	15	0	1470946865906	/Users/nikhilshekhar/h2o_working_directory/test/s3data_new/prostate_train.lib	gbm	gbm
2	5	0.402218	0.413716	0.324860	0.332081	0.491422	0.525792	0.223799	0.240132	0.834097	0.810320	15	0	1470946865906	/Users/nikhilshekhar/h2o_working_directory/test/s3data_new/prostate_train.lib	gbm	gbm
4	1	1.155843	76.990936	0.458816	57.607040	-6051.563477	-6010.870117	-164.260040	-163.155487	NULL	NULL	8	0	1470946866606	/Users/nikhilshekhar/h2o_working_directory/test/s3data_new/cars_train.lib	glm	glm
4	2	1.150747	76.994400	0.456809	57.607822	-6051.563477	-6010.870117	-164.260040	-163.155487	NULL	NULL	8	0	1470946866606	/Users/nikhilshekhar/h2o_working_directory/test/s3data_new/cars_train.lib	glm	glm
4	3	1.145684	76.997871	0.454821	57.608604	-6051.563477	-6010.870117	-164.260040	-163.155487	NULL	NULL	8	0	1470946866606	/Users/nikhilshekhar/h2o_working_directory/test/s3data_new/cars_train.lib	glm	glm
4	4	1.140651	77.001343	0.452839	57.609375	-6051.563477	-6010.870117	-164.260040	-163.155487	NULL	NULL	8	0	1470946866606	/Users/nikhilshekhar/h2o_working_directory/test/s3data_new/cars_train.lib	glm	glm
4	5	1.135651	77.004807	0.450869	57.610149	-6051.563477	-6010.870117	-164.260040	-163.155487	NULL	NULL	8	0	1470946866606	/Users/nikhilshekhar/h2o_working_directory/test/s3data_new/cars_train.lib	glm	glm
5	1	7.475157	7.574760	4.887924	4.941753	-789.414856	-791.454346	-21.432821	-21.482767	NULL	NULL	10	0	1470946869118	/Users/nikhilshekhar/h2o_working_directory/test/s3data_new/housing_train.lib	glm	glm
5	2	6.901058	7.000740	4.499975	4.571213	-789.762329	-791.454346	-21.440079	-21.482767	NULL	NULL	10	0	1470946869118	/Users/nikhilshekhar/h2o_working_directory/test/s3data_new/housing_train.lib	glm	glm
5	3	6.506752	6.598309	4.275152	4.344232	-792.080139	-791.454346	-21.499920	-21.482767	NULL	NULL	10	0	1470946869118	/Users/nikhilshekhar/h2o_working_directory/test/s3data_new/housing_train.lib	glm	glm
5	4	6.225984	6.312995	4.145546	4.211830	-793.167786	-793.388489	-21.529274	-21.544903	NULL	NULL	10	0	1470946869118	/Users/nikhilshekhar/h2o_working_directory/test/s3data_new/housing_train.lib	glm	glm
5	5	6.021369	6.107774	4.071876	4.132611	-793.411804	-795.514038	-21.529274	-21.592960	NULL	NULL	10	0	1470946869118	/Users/nikhilshekhar/h2o_working_directory/test/s3data_new/housing_train.lib	glm	glm
6	1	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	76934	0	1470946870537	/Users/nikhilshekhar/h2o_working_directory/test/s3data_new/mnist_train.lib	classifier	classifier
6	2	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	NULL	76934	0	1470946870537	/Users/nikhilshekhar/h2o_working_directory/test/s3data_new/mnist_train.lib	classifier	classifier

Learning outcomes

- Deep exposure to h2o, Spark ML and xgboost machine learning libraries.
- Stumbled upon multiple unresolved JIRA’s for spark and had to look for alternative solution to solve the problem at hand.
- Exposure to multiple datasets stored in different formats.
- Dealing with missing labels,fields in Spark-ml,xgboost,scikit.
- Reading data from s3 buckets programmatically using Java and Python.

Work in progress

- Run Spark in multi-node environment (Our own machines to start with, but EC2 if we ever decide to publish this framework)
- Run xgboost in a distributed environment with spark (Code is available in “ml-benchmark” repo, but needs testing)
- Tune the hyper-parameters for optimal performance of algorithms
- Add more test cases to each of the libraries (Utilize more “Kaggle like” datasets. Right now we have whatever is available in S3 for H2O)
- Run algorithms on large datasets (Besides the airlines dataset…)
- Integrating h2o accuracy suite - First, need to publish the test classes to maven central repository.
    - Work around is to paste in necessary classes into the “ml-benchmark” repo for H2O
- Benchmark.ai ? -> A combination of “ml-benchmark” and “db-benchmark”
- Other frameworks to benchmark?

Acknowledgements

- Navdeep Gill
- Bill Gallmeister
- Anmol Bal
- Nancy Jordan

Thank You